

ANUPREET SINGH

anupreet2226579@gmail.com | [Portfolio](#) | [LinkedIn](#) | [GitHub](#) | +1 (410) 260-9818

EDUCATION

University of Maryland, Baltimore County

(Master of Science in Computer Science, GPA: 3.8/4.0)

Baltimore, MD

Aug 2024 – May 2026

Indian Institute of Information Technology, Sonapat

(Bachelor of Technology in Computer Science and Engineering, GPA: 8.45/10)

Haryana, India

Aug 2019 – May 2023

SKILLS / RELEVANT COURSEWORK

Programming Languages: Python, C++, Javascript, SQL

Tools and Frameworks: FastAPI, Flask, sklearn, PyTorch, LangChain, Pincecone, MS SQL Server, PostgreSQL, React, AWS, Azure, Redis, Kafka, Docker, Kubernetes, GitHub Actions(CI/CD), Git, Linux

Areas of Interest and Coursework: AI & LLMs, Distributed System Design, Software Engineering (SDLC)

INDUSTRY EXPERIENCE

Data AI Engineer | Platinum Business Services (R&D Team) | Catonsville, MD

Dec 2025 – Present

- Developed **backend services** for an **AI-native** GRC platform that audits an SSP draft for compliance with NIST 800-53, chunking ~1,200 controls into a per-model vector index with linkage metadata tying each chunk to its parent control and sibling clauses, cutting a control-by-control review from **days of manual work to <10 minutes**.
- Built a priority-ordered fallback in the **RAG chat agent** across 5 embedding models (OpenAI, Anthropic, Ollama), re-embedding and re-running low-citation queries against each successive model's vector index to exploit differences in embedding representation and surface SSP evidence missed by higher-priority indexes, raising evidence-backed control assessment **accuracy from 64% to 92%**.
- Containerized the platform (FastAPI, Postgres, Qdrant, Ollama) and deployed it to **AWS EKS** with Kubernetes manifests and a GitHub Actions **CI/CD pipeline**, adding health checks and admin endpoints to download models and rebuild search indexes, so the team could manage models and indexes without shelling into pods.

Software Engineer (AI) Intern | PMA | Tallahassee, FL

Sep 2025 – Dec 2025

- Built an **AI-powered** PRD review platform in React and FastAPI that compares each uploaded draft PDF against a reference PRD and emits schema-validated gap flags for a **human-in-the-loop** reviewer to confirm or dismiss, giving a cohort-wide matrix that replaces a read of every draft with **LLM as judge** report and cuts feedback-session prep from **8 to 2 hours**.
- Architected a fully serverless review pipeline on AWS (S3 + CloudFront-hosted React SPA, Django on Lambda, DynamoDB, and an SQS-triggered Lambda AI worker) that retries failed analyses off a dead-letter queue and logs per-report model API spend, sustaining **200+** submissions across 5 cohorts at **~\$3/month** idle cost.

Software Development Engineer | Turfco Landscape Co. | Surrey, BC

Jan 2023 – June 2024

- Designed and deployed the **AWS** architecture that took the business off a WordPress brochure site, with SPA on CloudFront, autoscaling Express tier behind an ALB across two AZs, Multi-AZ RDS, sized for ~10x seasonal peaks and load-tested to **~50 quote requests/sec at <120 ms p99** with zero-downtime rolling deploys.
- Architected a self-serve quoting platform that priced jobs with deterministic multi-stage evaluation, JSONB encoded conditions, per-quote audit trails of every rule applied, and a write-invalidated Redis cache, all driven by rules editable from an admin console instead of a developer deploy, cutting quote turnaround from **~2 days to <1 min**.
- Built the online ordering and booking platform by taking self-serve orders through Stripe Checkout deposits, confirming them with signature-verified idempotent webhooks, and relaying a transactional outbox to Kafka workers for order confirmation, customer notification, and crew allocation, taking **\$600K+** in online bookings across ~1,100 orders.
- Delivered the React SPA that carried a customer from property details to instant estimate to paid deposit in one client-routed flow, with role-gated admin routes and JWT access tokens backed by httpOnly refresh cookies, enabling a streamlined booking flow with **~40%** of bookings occurring outside office hours without phone support.

PROJECTS

When Pixels Talk Back | Watermarking EMR Data to analyse impact on CV models

- Fine-tuned **3 computer-vision models** (ResNet50, DenseNet169, MobileNetV2) for classification over COVID-19 lung CT scans via two-stage ImageNet transfer learning, to establish a baseline peaking at **93%** accuracy and **0.951 F1**.
- Embedded a 1024-bit blind fragile watermark into the Region-of-Non-Interest of CT scans via key-scrambled LSB embedding of a SHA-512 fingerprint, enabling content authentication with near-lossless quality at **SSIM 0.999**.
- Benchmarked 3 watermarking algorithms across all 3 models, showing the fragile scheme costs **0%** accuracy loss, versus ~1% for reversible embedding with a 300-bit payload and 11.2% for robust embedding.